

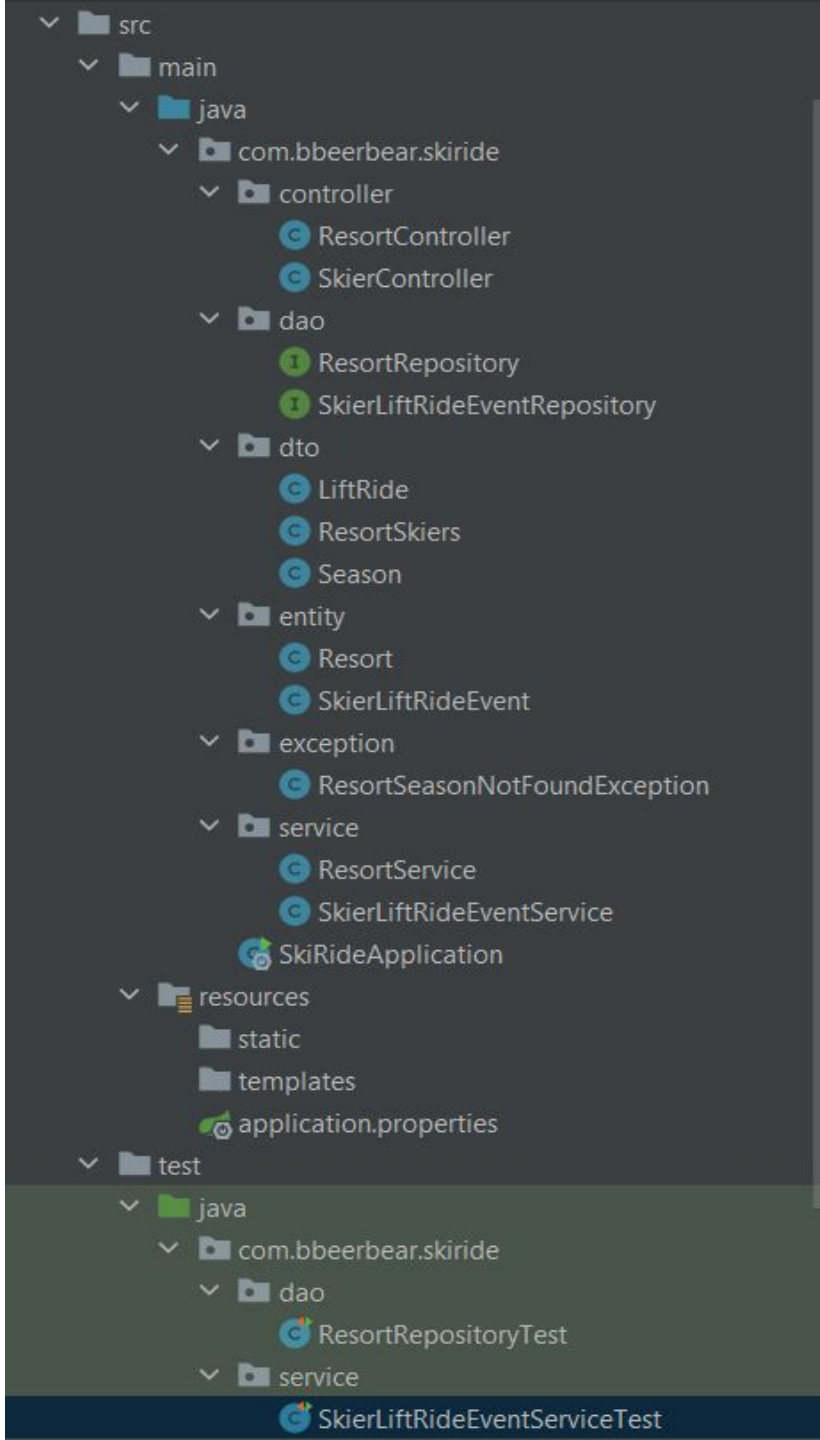
Digitalizing the Ski Sport Industry with a Cloud-Based Distributed System

Team member: Limin Xiong
Pranav Jha
Sahibmeet Singh
Vishruth Khatri

Presenter: Limin Xiong, Pranav Jha, Sahibmeet, Vishruth

Technology

- Server
 - Spring Boot
 - REST API
 - DTO
 - MongoDB
 - Oracle Cloud Deployment
- Client
 - Thread Concurrency



Server-side

Spring Data MongoDB

- MongoDB Template and MongoDB Repository
- Provides predefined methods

Spring Web Module

- Handling HTTP requests and responses, handling exceptions

Spring Test Module

- Provides a set of easy-to-use annotations and utility classes

Deployment

- On Oracle Cloud

Realize REST API and Save data to MongoDB

```
spring.data.mongodb.uri=mongodb+srv://bigbea...cluster1.pvadida.mongodb.net/?retryWrites=true&w=majority
spring.data.mongodb.database=skier_rider

3 usages  Limin
public interface SkierLiftRideEventRepository extends MongoRepository<SkierLiftRideEvent, String> {

}

@RestController
public class SkierController {
    2 usages
    private final SkierLiftRideEventRepository skierLiftRideEventRepository;
    2 usages
    private final SkierLiftRideEventService skierLiftRideEventService;
    no usages  Limin
    public SkierController(SkierLiftRideEventRepository skierLiftRideEventRepository, SkierLiftRideEventService skierLift
        this.skierLiftRideEventRepository = skierLiftRideEventRepository;
        this.skierLiftRideEventService = skierLiftRideEventService;
    }

    // Write a new lift ride for the skier
    no usages  Limin
    @PostMapping("/skiers/{resortID}/seasons/{seasonID}/days/{dayID}/skiers/{skierID}")
    public void createSkierLiftRide(@PathVariable int resortID, @PathVariable String seasonID,
        @PathVariable String dayID, @PathVariable int skierID,
        @RequestBody LiftRide liftRide){
        LiftRide liftRide1 = new LiftRide(liftRide.getLiftId(), liftRide.getTime());
        SkierLiftRideEvent skierLiftRideEvent = new SkierLiftRideEvent(id: null, resortID, seasonID, dayID, skierID, liftR
        skierLiftRideEventRepository.save(skierLiftRideEvent);
    }
}
```

Lift Ride Event Generator Thread

- **Store** the generating lift ride event in the **Blocking Queue** (thread-safe)
- Run the Thread **separately** and the client take the data from the queue.

```
// Create a blocking queue for lift ride events
```

```
2 usages
```

```
private final BlockingQueue<LiftRideEvent> queue = new LinkedBlockingQueue<>();
```

```
// Start the lift ride event generator thread
```

```
LiftRideEventGeneratorThread generatorThread = new LiftRideEventGeneratorThread(queue, num_of_thread,  
    num_of_requests_each_thread);
```

```
generatorThread.start();
```

```
LiftRideEvent liftRideEvent = queue.take();
```


Thread Concurrency

- Run threads with ExecutorService

```
for(int i = 0; i < num_of_thread; i++) {
    executorService.submit(()->{
        for(int j = 0; j < num_of_requests_each_thread; j++) {
            try {
                long requestStartTime = System.currentTimeMillis();
                LiftRideEvent liftRideEvent = queue.take();
                String url = "http://localhost:8080/coen6731/skiers/" + Integer.toString(liftRideEvent.getResortID()) +
                    liftRideEvent.getSeasonID() + "/days/" + liftRideEvent.getDayID() + "/skiers/" + Integer.toString(liftRideEvent.getDayID());
                String url = "http://155.248.230.86:8080/skiers/" + Integer.toString(liftRideEvent.getResortID()) + "/seasons/" +
                    liftRideEvent.getSeasonId() + "/days/" + liftRideEvent.getDayID() + "/skiers/" + Integer.toString(liftRideEvent.getDayID());
                String requestBody = new Gson().toJson(liftRideEvent.getLiftRide());

                HttpRequest request = HttpRequest.newBuilder()
                    .header("Content-Type", "application/json")
                    .uri(URI.create(url))
                    .POST(HttpRequest.BodyPublishers.ofString(requestBody))
                    .build();

                HttpResponse<String> response = client.send(request, HttpResponse.BodyHandlers.ofString());
                System.out.println(System.currentTimeMillis() - requestStartTime);
                eachRequestTimes.add(System.currentTimeMillis() - requestStartTime);

                int retries = 0;
                // Retry up to MAX_RETRIES times for 4XX and 5XX response codes
                while (response.statusCode() >= 400 && retries < MAX_RETRIES) {
                    retries++;
                    System.out.println("Request failed with status code " + response.statusCode() + ", retrying (attempt " + retries + ")");
                    response = client.send(request, HttpResponse.BodyHandlers.ofString());
                }
            } catch (Exception e) {
                e.printStackTrace();
            }
        }
    });
}
```

Terminate all threads

- Shutdown the ExecutorService gracefully
 - `executorService.shutdown();`
- Wait for all threads to complete
 - `executorService.awaitTermination(10, TimeUnit.MINUTES);`

```
try {  
    generatorThread.join();  
    executorService.shutdown();  
    executorService.awaitTermination(timeout: 10, TimeUnit.MINUTES);  
} catch (InterruptedException e1) {  
    // TODO Auto-generated catch block  
    e1.printStackTrace();  
}
```


Client 1

You have connectivity to call the API...

```
Number of threads used: 1000
Each request latency: 9389.11001905526ms
Number of successful requests sent: 9971
Number of unsuccessful requests: 29
Total run time: 99881 ms
Total throughput: 100.11914177871667 requests/second
According to Little's Law, expected throughput is: 106.5063672670246 requests/second

long endTime = System.currentTimeMillis();
long wallTime = endTime - startTime;
int successfulRequests = 0;
int unsuccessfulRequests = 0;
double throughput = 0;

int totalRequestTime = 0;
for(long time: eachRequestLatencies) {
    totalRequestTime += time;
}

System.out.println("Number of threads used: " + num_of_threads);
// expected throughput
double w_averageTimeEachRequest = totalRequestTime / successfulRequests;
double lambda_expectedThroughput = 1 / w_averageTimeEachRequest;

System.out.println("Each request latency: " + w_averageTimeEachRequest);
System.out.println("Number of successful requests sent: " + successfulRequests);
System.out.println("Number of unsuccessful requests: " + unsuccessfulRequests);
System.out.println("Total run time: " + wallTime + " ms");
System.out.println("Total throughput: " + throughput + " requests/second");
System.out.println("According to Little's Law, expected throughput is: " + lambda_expectedThroughput + " requests/second");
```


Client2 - Output - Profiling Performance

```
double meanResponseTime = (double) totalResponseTime / eachRequestTimes.size();
double medianResponseTime = MedianCalculator.calculateMedian(eachRequestTimes);

long endTime = System.currentTimeMillis();
long wallTime = endTime - startTime;
double throughput = (double) num_of_thread * num_of_requests_each_thread / (wallTime / 1000.0);

long p99ResponseTime = P99Calculator.calculateP99(eachRequestTimes);

long minResponseTime = Collections.min(eachRequestTimes);
long maxResponseTime = Collections.max(eachRequestTimes);

// expected throughput  $\lambda = L / W$ 
double w_averageTimeEachRequest = (double) totalResponseTime / num_of_requests_each_thread;
double  $\lambda$ _expectedThroughput = num_of_thread / (w_averageTimeEachRequest / 1000.0);

System.out.println("Number of threads used: " + num_of_thread);
System.out.println("Number of requests each thread: " + num_of_requests_each_thread);
System.out.println("Mean response time: " + meanResponseTime + "ms");
System.out.println("Median response time: " + medianResponseTime + "ms");
System.out.println("Throughput: " + throughput + " requests/second");
System.out.println("According to Little's Law, expected throughput is: " +  $\lambda$ _expectedThroughput + " requests/second");
System.out.println("P99 response time: " + p99ResponseTime + "ms");
System.out.println("Min response time: " + minResponseTime + "ms");
System.out.println("Max response time: " + maxResponseTime + "ms");
```

Mean response time: 994.8844ms
Median response time: 1001.0ms
Throughput: 99.91207737191273 requests/second
According to Little's Law, expected throughput is: 100.51419039237122 requests/second
P99 response time: 1076ms
Min response time: 120ms
Max response time: 1231ms

Little's Law

- According to Little's Law: $L = \lambda W$
- λ = throughput
- L = total number of requests
- W = average time of each request
- In 200 threads, each with 50 requests,
 - The total throughput is **99.75360858679063** requests/second
 - Expected throughput is **101.365335315769** requests/second

```
Number or threads used: 200, number of requests each thread: 50
```

```
Each request latency: 1973.0611ms
```

```
Number of successful requests sent: 10000
```

```
Number of unsuccessful requests: 0
```

```
Total run time: 100247 ms
```

```
Total throughput: 99.75360858679063 requests/second
```

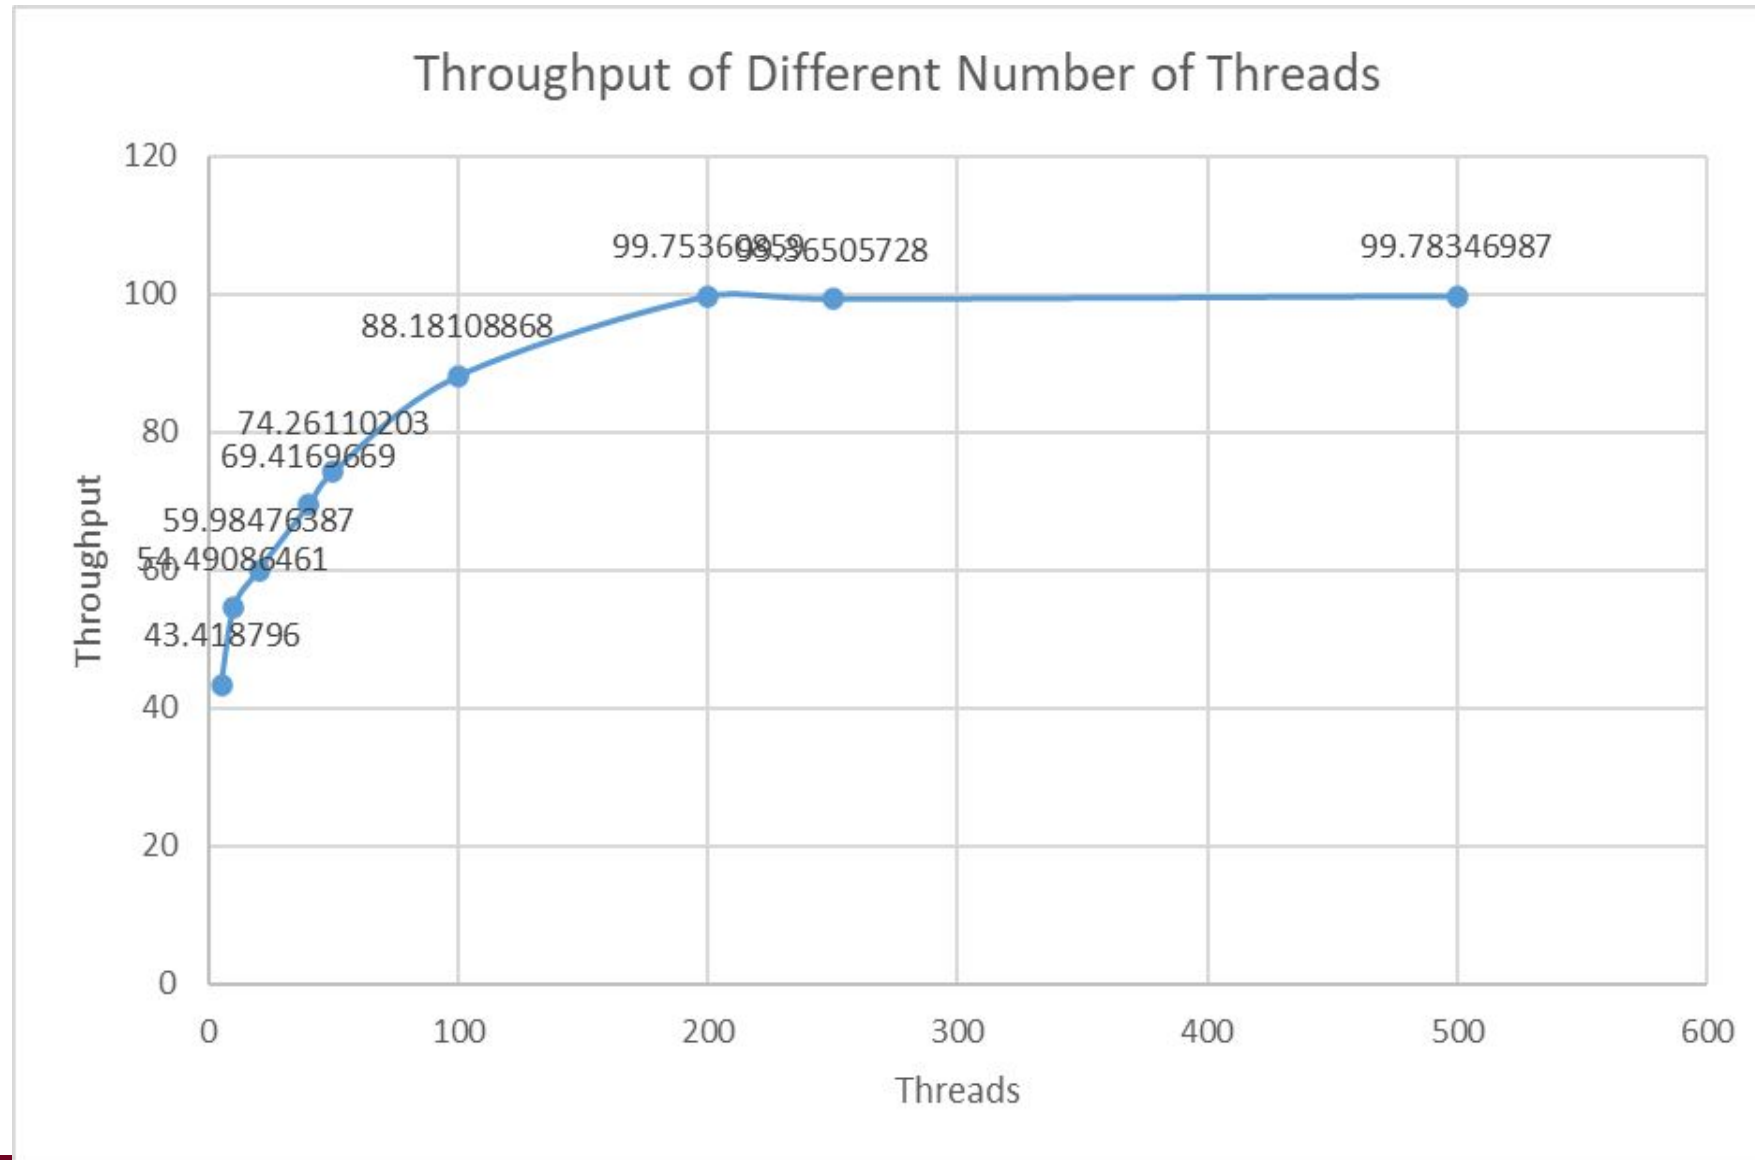
```
According to Little's Law, expected throughput is: 101.365335315769 requests/second
```

Maximum of the throughput(10k requests)

Testing the throughput in different number of threads

Threads	Throughput(requests/s)
5	43.418796
10	54.49086461
20	59.98476387
40	69.4169669
50	74.26110203
100	88.18108868
200	99.75360859
250	99.36505728
500	99.78346987

Because of the thread concurrency overhead, the **throughput** is **bounded** by the capacity



API Swaggerhub

skier-controller

POST /skiers/{resortID}/seasons/{seasonID}/days/{dayID}/skiers/{skierID}

Parameters Try it out

Name	Description
resortID * required integer(\$int32) (path)	<input type="text" value="resortID"/>
seasonID * required string (path)	<input type="text" value="seasonID"/>
dayID * required string (path)	<input type="text" value="dayID"/>
skierID * required integer(\$int32) (path)	<input type="text" value="skierID"/>

Request body required application/json

POST /resorts

Parameters Try it out

No parameters

Request body required application/json

Example Value | Schema

```
{  "resortID": 0,  "resortName": "string",  "seasons": [    "string"  ]}
```

Responses

Code	Description	Links
200	OK	No links

Media type */

Controls Accept header.

Example Value | Schema

```
{  "resortID": 0,  "resortName": "string",  "seasons": [    "string"  ]}
```

Link for the Swaggerhub API -:

<http://155.248.230.86:8080/swagger-ui/index.html#/>